

NEXTELEVEN LLC

PARTS

Multi-Location Dealership Parts Intelligence

Customer Sales Packet
Proprietary Product Information
July 2026

One sentence: *Parts finds the right part across every rooftop in your group — in plain English — and tells you green, yellow, or red whether a human should touch the answer.*

This packet is for dealership groups evaluating Parts for multi-store availability, counter productivity, and trustworthy AI at the parts desk. It is informational and does not grant a software license. All software remains proprietary to NextEleven LLC.

Fact	Detail
Product	Parts — Multi-Location Dealership Parts AI
Company	NextEleven LLC (Frisco, Texas)
Engineering reference	https://github.com/seanebones-lang/parts
Core package	parrts v0.4.2 (hybrid RAG)
License	Proprietary — not open source
Last measured health	2026-07-31 · 59 tests passed · eval 12/12 hit@5
Pilot shape	30 days · 1–3 rooftops · written success metrics

Contents

- 1. The Problem We Solve
- 2. Solution Overview
- 3. How It Works (Architecture)
- 4. Key Capabilities
- 5. What Is Real vs Optional
- 6. Traffic-Light Trust Model
- 7. Demo Scenario and Proof Points

- **8.** Ideal Fit and Disqualifiers
- **9.** 30-Day Pilot Blueprint
- **10.** ROI Discussion Framework
- **11.** Security, Data and Licensing
- **12.** Implementation and Requirements
- **13.** FAQ
- **14.** Commercial Next Steps
- **15.** Appendix: Technical Health Snapshot

1. The Problem We Solve

Across multi-rooftop dealer groups, a simple question — “*Do you have brake pads for a 2019 Honda Civic?*” — still burns minutes of advisor and counter time. Staff bounce between DMS screens, phone trees, and tribal knowledge. Wrong “yes” answers create comebacks, goodwill costs, and CSI risk. Wrong “no” answers lose work to competitors.

Generic chatbots make this worse when they invent inventory. Dealers do not need more confident paragraphs. They need **defensible availability answers** with a clear human gate.

Typical pain points

- Cross-store blindness — nobody sees group inventory in one plain-English step
- Advisor hold time — every availability check steals fixed-ops capacity
- Inconsistent process — each rooftop answers differently
- AI distrust — prior chatbot pilots hallucinated stock
- Integration fear — teams reject rip-and-replace DMS projects

2. Solution Overview

Parts is multi-location dealership parts intelligence from NextEleven LLC. Users ask in natural language. The system retrieves and ranks parts across locations, applies a green / yellow / red traffic-light policy, and exposes results via CLI, API, and web UI.

Parts is designed to sit **beside** your DMS — accelerating availability intelligence — not to rip out your system of record on day one.

Value pillars

Pillar	What the customer gets
Find	Plain-English lookup that understands vehicle + part language
Trust	Traffic-light policy so the floor knows when AI is enough
Multi-store	Ranked results across rooftops; expand sibling locations by SKU
Integrate later	FastAPI enterprise surface when you are ready for DMS/SSO
Measure	Golden-query eval harness for pilot proof — not vibes

Positioning statement

“Parts is retrieval-first dealership software: plain-English multi-store availability with a human brake pedal — not chatbot theater and not a DMS replacement.”

3. How It Works (Architecture)

Parts is a dual-stack platform. The **parrts core** is the source of truth for retrieval. The enterprise backend and frontend sit on top for APIs, agents, and UI.

Layer	Role	Customer meaning
parrts core	Hybrid dense + keyword ranking, traffic-light policy, offline-capable	The brain for “do you have it?”
Enterprise API	FastAPI, /api/v1 modules, health, metrics	Integration surface for IT
LangGraph agents	Parts, CS, inventory, pricing, payment, shipping, supplier, follow-up	Workflow orchestration; soft-fail offline
Frontend	Next.js parts search + traffic badge	Counter / pilot UI
Ops	Docker profiles, CI, eval, pgvector seed	Deploy path when ready

Default embedding path for demos is offline-safe (no cloud model required). Optional cloud LLMs can help phrase answers but are not required for stock retrieval. Optional Postgres + pgvector supports larger centralized catalogs when available.

4. Key Capabilities

Natural language parts lookup. Ask the way humans ask — vehicle, part type, brand cues.

Hybrid retrieval. Combines meaning-based and keyword-style ranking (RRF fusion) for parts language.

Multi-location ranking. See which rooftop can fulfill; filter by location when needed.

Parent / sibling expand. Expand a hit across stores carrying the same base SKU.

Traffic-light policy. Green / yellow / red with confidence and recommended actions.

CLI for power users. parrts ingest / query / status — scriptable and demo-friendly.

HTTP API. /query, /health, /metrics plus enterprise /api/v1 route groups.

Web UI. Parts page with live API path and mock fallback for resilient demos.

Agent orchestration. Email to classify to parts/service/inventory/pricing paths with human flags.

Evaluation harness. Fixed golden queries; hit@k, MRR, recall, latency reporting.

Auth modes. AUTH_MODE=demo for controlled demos; production expects JWT on wired writes.

pgvector path. Opt-in production vector store with seed script when Postgres is up.

5. What Is Real vs Optional (Honesty Table)

We win long dealer relationships by not overselling. Use this table in every serious conversation.

Capability	Status today	What you need
Hybrid RAG lookup + traffic-light	Measured / shipping	Install core package
Multi-location seed catalog demo	Ready	Laptop or server
Automated tests + eval harness	Measured	None
FastAPI boot + 15 API modules load	Measured	Python deps
Next.js UI type-check clean	Measured	Node for UI
Cloud LLM answer wording	Optional	Anthropic/OpenAI keys
Live Stripe payments	Mock until keyed	Stripe account + scope
Live carrier labels (EasyPost)	Mock until keyed	EasyPost + scope
Live pgvector at scale	Code-ready	Docker + Postgres
Full DMS bi-directional sync	Integration SOW	Discovery + project
Unattended full autonomy / no humans	Not offered as SLA	—
Open-source / free redistribution	Not available	Proprietary license only

Claims of full 13-agent production SLA, guaranteed ROI dollars, or always-live supplier scraping are aspirational unless scoped in a signed SOW.

6. Traffic-Light Trust Model

Traffic-light is a product feature for fixed-ops culture — not a slide gimmick.

Signal	Typical meaning	Recommended floor behavior
GREEN	Strong match and healthy stock posture	Safe to answer quickly; still verify process rules
YELLOW	Mid confidence and/or thin stock	Human reviews before promising customer
RED	Weak match, zero stock, or no hits	Do not auto-promise; escalate / source

Thresholds can be tuned in a pilot against your golden queries and risk tolerance. The goal is not zero human involvement — it is putting humans on the hard cases, not every easy availability check.

7. Demo Scenario and Proof Points

Example demo query

parrts query "brake pads for 2019 Honda Civic" --no-llm

On the standard seed catalog (health check 2026-07-31): traffic-light **green**, top SKU example **BP-HC19-L4**, multi-location hits available; --expand-parent shows sibling rooftops for the same base part.

Measured engineering proof (as of 2026-07-31)

Check	Result
Automated tests	59 passed, 2 skipped (optional ST + live DB by design)
Retrieval eval (default set)	hit@5 = 12/12 (100%), MRR = 1.0, recall@5 = 1.0
Mean eval latency (demo hardware)	About 0.7 ms per query on core path
API v1 modules loaded	15 / 15 (0 failed)
Core endpoints without Postgres	/, /health, /query, /metrics available
Frontend TypeScript	tsc --noEmit clean
Compose config	Valid profiles: core, api, full, pgvector, obs

Your pilot replaces the seed catalog with your data and rebuilds a golden query set you own.

8. Ideal Fit and Disqualifiers

Strong fit

- Multi-rooftop groups (roughly 3–15 stores) with shared or loosely shared parts inventory
- Parts managers and fixed-ops leaders under pressure on speed-to-answer and CSI
- Groups burned by chatbots that invented stock
- Willingness to run a measured 30-day pilot with a golden query list

Usually a weak fit (for now)

- Single-store shops that only need basic e-commerce catalog search
- Buyers demanding fully autonomous ordering with zero humans on day one
- Prospects unwilling to measure quality or name a business owner
- Requests for open-source redistribution without agreement

9. 30-Day Pilot Blueprint

Week	Activities	Exit artifact
0	Kickoff, NDA if needed, rooftop list, owner names	Charter + success metrics
1	Install core; index seed or sample export; draft golden queries	Baseline hit-rate
2	Shadow mode with counter/advisors; capture overrides	Usage notes
3	Tune traffic-light; optional LLM wording; train cohort	Updated thresholds
4	Readout: quality, latency, overrides, go/no-go	Decision + production plan

Suggested success criteria (negotiate per group)

- hit@5 at least 90% on the agreed golden query set
- Median latency under 2 seconds on agreed hardware for interactive use
- Yellow/red rate understood and staffable
- Counter cohort would use daily (simple survey)

10. ROI Discussion Framework

The following is a **discussion framework**, not a guarantee. Replace every input with dealer actuals.

Status-quo cost drivers

Driver	Example input	Why it matters
Daily “do you have” events (group)	200	Volume of friction
Minutes per event today	4	Labor sink
Loaded labor \$/hour	\$35	Converts time to money
Wrong-promise / comeback rate	3%	Hidden rework + CSI
Cost per comeback	\$75	Goodwill + rework

Illustrative only: 200 events times 4 minutes is about 13.3 hours/day — on the order of six figures/year in lookup labor before comeback costs. A conservative pilot that improves a fraction of events by a couple of minutes, plus a small drop in wrong promises, is often enough to justify a pilot fee — **if measured on your data**.

We sell the pilot on measurement: hit-rate, latency, override rate, and a before/after time sample — not on a magical ROI spreadsheet.

11. Security, Data and Licensing

- **Proprietary IP.** Parts is not open source. Copyright and trade-secret protections apply.
- **No implied license** from viewing this packet or a repository reference URL.
- **Auth modes:** demo for controlled demos; production mode expects JWT on wired mutating routes.
- **Network:** do not expose admin APIs publicly without authentication and controls.
- **Processors:** Stripe / EasyPost / LLM providers only when you supply keys and approve scope.
- **Customer data:** under NDA/MSA; no training unrelated public models on your catalog without agreement.
- **Deployment:** laptop pilot, on-prem/VPC, or controlled cloud — scoped in enterprise SOW.

12. Implementation and Requirements

Technical prerequisites (pilot)

- Python 3.11+ environment (venv recommended)
- Optional: Node 20+ for web UI
- Optional: Docker Desktop for Postgres/pgvector full stack
- Named business owner (parts or fixed ops) + technical contact
- Location list and naming conventions
- Top 20–50 real availability questions (golden set)
- Sample catalog extract or agreement to start on demo catalog then swap

Quick start (technical evaluation)

```
git clone https://github.com/seanebones-lang/parts.git
cd parts
python3 -m venv .venv
source .venv/bin/activate
pip install -e ".[dev]"
python -m parrts ingest --force
python -m parrts query "brake pads for 2019 Honda Civic" --no-llm
pytest -q
python scripts/eval_retrieval.py
```

13. FAQ

Q: Does this replace our DMS?

A: No. It accelerates availability intelligence and can integrate over time via API.

Q: Will it invent stock levels?

A: Core design is retrieval-first with traffic-light gating. Weak results surface as yellow/red for humans.

Q: Do we need cloud AI keys?

A: Not for core lookup demos. Keys unlock optional answer wording and live processors.

Q: How fast to first value?

A: Same-day demo on sample catalog; pilot value measured in weeks.

Q: Can it run in our VPC?

A: Yes — architecture supports controlled deployment; scope in enterprise SOW.

Q: Is the GitHub repo free to use?

A: No. Proprietary. Access and use require NextEleven authorization.

Q: What about payments and shipping?

A: Mocked unless you connect Stripe/EasyPost and contract that scope.

Q: Who owns the golden queries?

A: You do. We help facilitate; parts leadership should own the list.

14. Commercial Next Steps

Recommended path

- **Step 1.** Discovery call — rooftops, DMS, pain, owners (30 minutes)
- **Step 2.** NDA if required by your legal team
- **Step 3.** Live demo on your top questions (or seed catalog if export pending)
- **Step 4.** Pilot SOW — 30 days, 1–3 rooftops, written metrics
- **Step 5.** Readout and production decision (group license + optional integration SOW)

Commercial framing (ranges set in formal quote)

Motion	What it covers
Pilot	Time-boxed fee, success criteria, limited rooftops
Group subscription	Annual license by rooftop band + support
Enterprise integration	DMS/SSO/VPC, SLA, training, custom feeds

Contact your NextEleven LLC representative to schedule discovery.

Engineering reference: <https://github.com/seanebones-lang/parts>

Company: NextEleven LLC · Frisco, Texas · Proprietary and confidential

Appendix A — Technical Health Snapshot (2026-07-31)

Internal automated health check summary for sales engineering alignment. Not a third-party audit.

Area	Status	Notes
Core hybrid RAG	PASS	Offline retrieval + traffic-light
Pytest	PASS	59 passed / 2 skipped
Backend /api/v1 boot	PASS	15/15 routers
Frontend tsc	PASS	exit 0
Retrieval eval	PASS	12/12 hit@5, MRR 1.0
Docker live stack	ENV BLOCKED	Daemon off at check time; compose YAML valid
pgvector seed	GRACEFUL SKIP	Exit 0 when DB unreachable
Live LLM/Stripe/EasyPost	N/A	Key-gated by design
Ruff style	WARN	About 24 non-blocking style findings

Appendix B — Sample Interaction (Illustrative)

Input: brake pads for 2019 Honda Civic

Output (conceptual): ranked SKUs with location, stock, scores; traffic-light color; optional expand to other stores; optional natural-language summary if LLM keys present.

Appendix C — Document Control

Field	Value
Document	Parts Customer Sales Packet (PDF)
Version	1.0
Date	2026-07-31
Audience	Prospective dealer-group buyers and evaluators
Owner	NextEleven LLC
Classification	Confidential — customer evaluation

(c) 2026 NextEleven LLC. All rights reserved. This document does not grant any license to use, copy, modify, or distribute the Parts software or related intellectual property. Figures and pilot outcomes depend on customer data quality, configuration, and adoption. ROI examples are illustrative only.